

Machine Learning by Stanford University

Study Sheet by Reah Miyara ✉ ml@reah.me

Intro to Machine Learning


- **ML** – a computer program with increased performance P at some class of tasks T with experience E .
- **Supervised** – given a [‘ground truth’] data set, predict output given the input. Types of prediction:

1. **Regression** – continuous, numerical
2. **Classification** – discrete, categorical

- **Unsupervised** – derive structure from data based on relationships among variables (with no prior knowledge as to what the results should look like)

Linear Regression with One Variable

- **Learning Goal** – given a training set, learn a function $h: X \rightarrow Y$ so $h(x)$ is a good y predictor



- **Hypothesis** – $h_\theta(x) = \theta_0 + \theta_1 x$
- **Cost Function** – takes an average difference of all results of the hypothesis with inputs from the x values and the actual y values. Goal: minimize θ_0, θ_1

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_\theta(x_i) - y_i)^2 \quad (1)$$

(1) Squared Error function or Mean Squared Error function

- **Gradient Descent Algorithm** repeat until convergence

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (2)$$

Multivariate Linear Regression

$$h_\theta(x) = [\theta_0 \quad \theta_1 \quad \dots \quad \theta_n] \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix} = \theta^T x$$

- **Gradient Descent Algorithm** repeat until convergence

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)} ; j := 0 \dots n \quad (3)$$

- **Feature Scaling** – divide the input values by the range (max – min). Input values in roughly the same range speed up the convergence of gradient descent.
- **Mean Normalization** – subtract the mean for an input variable from the values for that input variable.

$$x_i := \frac{x_i - \mu_i}{s_i} \quad (4)$$

(4) μ_i is the mean & s_i is the range, (max – min), of all values for feature i

- **Learning Rate** – α too small \Rightarrow slow convergence; α too large \Rightarrow may not converge.

Normal Equation

- **Normal Equation** – non-iterative algorithm for minimizing $J(\theta)$; note : $O(n^3)$ to calculate $X^T X$

$$\theta = (X^T X)^{-1} X^T y \quad (5)$$

m examples $(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})$; n features

$$x^{(i)} = \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ \vdots \\ x_n^{(i)} \end{bmatrix} \in \mathbb{R}^{n+1} \quad | \quad X = \begin{bmatrix} (x^{(1)})^T \\ (x^{(2)})^T \\ \vdots \\ (x^{(n)})^T \end{bmatrix} \quad | \quad y = \begin{bmatrix} y^1 \\ y^2 \\ \vdots \\ y^m \end{bmatrix}$$

$x^{(i)}$ = training vector i (containing values from all features); $X \rightarrow m \times (n+1)$

- If $X^T X$ is noninvertible, common causes include:

1. Redundant features, where two features are very closely related (i.e. they are linearly dependent)
2. Too many features (e.g. $m \leq n$). In this case, delete some features or use ‘regularization’.

Classification Logistic Regression

- **Logistic Regression Model** – want $0 \leq h_\theta(x) \leq 1$

$$g(z) = \frac{1}{1 + e^{-z}} \quad (6)$$

- **Sigmoid/Logistic Function**

$$h_\theta(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}} \quad (7)$$

$$h_\theta(x) = P(y = 1|x; \theta) = 1 - P(y = 0|x; \theta) \quad (8)$$

- **Decision Boundary**

1. $y = 1$ if $\theta^T(x) \geq 0$
2. $y = 0$ if $\theta^T(x) < 0$

- **Cost Function**

$$Cost(h_\theta(x), y) = \begin{cases} -\log(h_\theta(x)) & \text{if } y = 1 \\ -\log(1 - h_\theta(x)) & \text{if } y = 0 \end{cases} \quad (9)$$

$$Cost(h_\theta(x), y) = -y \log(h_\theta(x)) - (1 - y) \log(1 - h_\theta(x)) \quad (10)$$

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m Cost(h_\theta(x^{(i)}), y^{(i)}) \quad (11)$$

a vectorized implementation is:

$$h = g(X\theta); \quad J(\theta) = \frac{1}{m} \cdot (-y^T \log(h) - (1 - y)^T \log(1 - h)) \quad (12)$$

$$\theta_j := \theta_j - \frac{\alpha}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)} \quad (13)$$

a vectorized implementation is:

$$\theta := \theta - \frac{\alpha}{m} X^T (g(X\theta) - \vec{y}) \quad (14)$$

Multiclass Classification

$y \in 0, 1, \dots, n$

$$h_\theta^{(0)}(x) = P(y=0|x; \theta)$$

$$h_\theta^{(1)}(x) = P(y=1|x; \theta)$$

\vdots

$$h_\theta^{(n)}(x) = P(y=n|x; \theta)$$

$$\text{prediction} = \max_i (h_\theta^{(i)}(x))$$

Under/Over –fitting

- **Underfitting** – hypothesis shows structure not captured by the model; does not fit the data well
- **Overfitting** – hypothesis corresponds too closely to data \therefore fail to predict future results reliably
- Addressing the problem of overfitting:
 1. **Reduce the number of features** – manually/algorithmically select subset of features.
 2. **Regularization** – keep all features, but reduce the magnitude of parameters θ_j by λ (regularization parameter).

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right] \quad (15)$$

Gradient Descent regularized linear regression does not penalize θ_0

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_0^{(i)} \quad (16)$$

$$\theta_j := \theta_j - \alpha \left[\left(\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)} \right) + \frac{\lambda}{m} \theta_j \right] \quad (17)$$

$$\theta_j := \theta_j \left(1 - \alpha \frac{\lambda}{m} \right) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)} \quad (18)$$

Normal Equation

$$\theta = (X^T X + \lambda \cdot L)^{-1} X^T y \quad (19)$$

where $L = \begin{bmatrix} 0 & & & \\ & 1 & & \\ & & \ddots & \\ & & & 1 \end{bmatrix}$